TurboSOAP: Compiled Dense Basis Tracking with BF16 State

Ethan Smith
ethansmith2000@gmail.com

Working review draft | 29 October 2026

Abstract

SOAP's matrix preconditioners introduce state and basis-refresh costs that can limit its practical efficiency. We describe TurboSOAP, a dense two-sided SOAP approximation combining damped basis tracking, periodic QR repair, BF16 persistent state, and compiled updates over equal-shaped matrices. Its precision policy retains FP32 master parameters and decomposition islands while using BF16 for moments, covariance factors, bases, and most matrix operands. On one 109-million-parameter language model trained on a fixed OpenWebText cohort, two model-initialization seeds finish 10,000 optimizer calls with cross-entropy increases of 0.01124 and 0.01008 relative to a stock-beta SOAP QR10 reference, corresponding to 1.13% and 1.01% higher perplexity. These endpoints follow larger observed gaps of 0.118–0.124 CE at call 500. Persistent matrix state falls from 2.320 to 1.160 GiB. Complete-step time falls 18.66% in a paired 9,900-call window for one seed and 17.36% in a matched 100-call prefix for the other. Per-event accounting places most optimizer-time saving in changed refresh work. Short matched-precision QR20/QR50 controls and a same-state QR profile provide cadence and implementation context. These measurements support a scoped quality/time/memory tradeoff; they do not establish statistical equivalence or attribute the combined speed improvement solely to basis tracking.

1. Introduction

SOAP applies Adam-style adaptation in the coordinate system of matrix preconditioners. Computing and maintaining that coordinate system adds work beyond diagonal optimizers. A practical approximation must address both its numerical history and how that history is stored and executed on a GPU.

We investigate one concrete implementation point: full dense factors and two-sided matrix coverage, BF16 persistent state, a fixed mixture of cheaper dense tracking steps and periodic QR refreshes, and compilation of the complete matrix-update regions. The objective is to document how closely this fixed implementation follows an audited SOAP reference on an affordable language modeling workload, together with the measured time and memory tradeoff.

The evidence has three parts. First, two initialization seeds compare the declared 10,000-call endpoint against stock-beta QR10. Second, short full-BF16 QR20 and QR50 controls test whether intermediate tracking provides useful learning behavior relative to simpler refresh schedules. Third, a fixed-input profile separates redundant products in QR refinement from the solver cost. The policy, precision rules, production optimizer hashes, and shared training-worker hash remain fixed across the two main candidate runs; campaign-specific runners are separately pinned.

Our contribution is the specified dense tracking/repair/precision implementation and its reproducible case study. The learning curves retain a measurable loss gap, especially early in training. The timing comparison includes storage, batching, arithmetic, compilation, and cadence changes relative to the original reference. These components should remain visible when interpreting efficiency.

2. Relationship to existing work

Vyas et al. describe SOAP's rotated-coordinate adaptive updates and the use of infrequent basis refinement: a warm-started power iteration followed by QR (Algorithm 4). Thus both arms already reuse the preceding basis. SOAP also discusses lower-precision state and the FP32 minimum for its QR implementation (§§7.2.1, 7.3.1, 8). We retain the FP32 solver island, use QR every 50 ages with intermediate dense motion, and store persistent

matrix state in BF16. Our stock matrix memory settings, model and auxiliary recipe are specified separately. SOAP, v2.

Milligan et al. store the basis-coordinate preconditioner and use subspace updates, including QR within selected blocks; they also study BF16 storage. Our implementation instead stores world-coordinate covariance factors and updates the full basis with dense damped motion and fixed QR repair. We do not maintain their persistent basis-coordinate representation or establish a performance ranking against it. Milligan et al., 2026, §§3.1–3.2.

Eschenhagen et al. use approximation-error criteria to adapt eigenbasis computation and examine the effects of stale bases. Our schedule is fixed across factors and training stages; no residual controls refresh decisions. Their analysis is relevant context, rather than a directly measured comparator in this study. Purifying Shampoo, v2, §4.

Warm reuse, subspace updates, low-precision storage, and compilation are existing ideas. The distinction studied here is their particular dense update, repair, state-history, and precision policy. Direct comparisons with subspace or adaptive implementations remain outside this experiment. Low-precision bases plus orthogonality correction are also prior work: Wang et al. quantize Shampoo eigenvector matrices and apply Björck orthonormalization (§3.2). Our BF16 dense SOAP policy is not a new instance of that general idea. 4-bit Shampoo, v3. The numerical ingredients also have classical lineage: Jacobi eigensolvers use plane rotations (Forsythe and Henrici, 1960), and iterative polar orthogonalization predates this work (Björck and Bowie, 1971).

3. Method

3.1. State, precision, and update ordering

For a matrix parameter $W \in \mathbb{R}^{m\times n}$, the optimizer stores first and second moments M, V, covariance factors C_L, C_R, and square bases Q_L, Q_R. All six persistent tensors use BF16. Parameters, scalar age and learning-rate tensors, pointwise moment temporaries, and decomposition inputs/outputs before basis storage use FP32. Auxiliary parameters use the separate AdamW optimizer described in Section 4.

Let R_b denote rounding to BF16. We distinguish two matrix products:

$$\mathcal{B}(X,Y) = R_b(XY), \qquad \mathcal{F}(X,Y) = \mathrm{FP32out}(XY),$$

where both operators receive BF16 operands and use FP32 accumulation. The notation specifies tensor rounding boundaries; it does not promise a single GPU kernel or the absence of other temporary buffers.

The reported CUDA runs disable BF16 reduced-precision reductions and set FP32 matmul precision to "high". FP32 denotes tensor and solver precision; FP32-tensor matrix products inherit that backend setting, which can use lower-precision internal multiplication. Geometry diagnostics temporarily use "highest". These settings apply to reference and candidate; casting QR inputs to FP32 does not request "highest" matmul precision. The harness also enables deterministic algorithms; the reproduction guide records the cuBLAS workspace, thread and compiler-cache settings.

On call $k = 1$, the matrix optimizer initializes zero moments, accumulates covariances from the clipped gradient, and obtains descending initial bases by FP32 eigendecomposition of symmetrized factors. It skips the matrix parameter update. Auxiliary AdamW still updates on that call. For subsequent calls define the matrix age $t = k - 1$, and let $G_b = R_b(G)$ for the current clipped gradient. In the existing bases,

$$\bar{G} = \mathcal{B}(\mathcal{B}(Q_L^\top, G_b), Q_R),$$

$$M_f = \beta_1 \,\mathrm{float}(M) + (1-\beta_1)\,\mathrm{float}(\bar{G}), \quad V_f = \beta_2 \,\mathrm{float}(V) + (1-\beta_2)\,\mathrm{float}(\bar{G})^2.$$

The current direction uses these unrounded FP32 moment temporaries:

$$U = \mathcal{F}\left(\mathcal{B}\left(Q_L, R_b\left[\frac{M_f}{\sqrt{V_f}+\epsilon}\right]\right), Q_R^\top\right), \qquad W^+ = \left(W - \eta_k \frac{\sqrt{1-\beta_2^t}}{1-\beta_1^t} U\right)(1-\eta_k\lambda).$$

The optimizer then stores $R_b(M_f), R_b(V_f)$, updates the covariances, and performs any scheduled basis change and moment transport. Thus the direction precedes this call's basis refresh. Rounding moments before using them for

the current direction would specify a different implementation.

The tensor-level covariance update is expressed as

$$C_L^+ = R_b\big[\beta_C \,\mathrm{float}(C_L) + (1-\beta_C) G_b G_b^\top\big], \quad C_R^+ = R_b\big[\beta_C \,\mathrm{float}(C_R) + (1-\beta_C) G_b^\top G_b\big].$$

The implementation uses BF16 `baddbmm` with FP32 accumulation for this operation. The equations specify the weighted update and its output dtype; they do not establish the vendor kernel's exact internal rounding sequence. All main experiments use $\beta_1 = \beta_2 = \beta_C = 0.95$.

3.2. Dense warm tracking

For either updated covariance factor C and its current basis Q, form (with $\mathrm{sym}(X) = (X + X^\top)/2$)

$$B = \mathrm{sym}\big(\mathcal{F}(\mathcal{B}(Q^\top, C), Q)\big), \quad d_i = B_{ii}, \quad s = \max(\max_i |d_i|, 10^{-30}), \quad h_{ij} = (d_i - d_j)/s.$$

We apply a damped skew motion in the estimated basis coordinates:

$$A_{ij} = -\frac{1}{2}\frac{B_{ij}}{s}\frac{h_{ij}}{h_{ij}^2 + 0.01^2}, \quad a = \min\left(1, \frac{0.5}{\max(\|A\|_F/\sqrt{2}, 10^{-12})}\right), \quad Y_0 = R_b\big(\mathrm{float}(Q) + \mathcal{F}(Q, R_b(aA))\big).$$

Two residual corrections follow:

$$E_j = \mathcal{F}(Y_j^\top, Y_j) - I, \qquad Y_{j+1} = R_b\big(\mathrm{float}(Y_j) - \tfrac{1}{2}\mathcal{F}(Y_j, R_b(E_j))\big), \quad j = 0, 1.$$

The stored new basis is Y_2. This is an approximate orthogonality correction, not exact QR. All products are dense: the motivation is GPU execution and numerical policy, not a reduction in cubic asymptotic order. The skew direction is a damped first-order Jacobi-type motion applied to all coordinate pairs. For a small skew A, the unrounded change is $B^+ \approx B + BA - AB$; its diagonal approximation gives the gap-dependent off-diagonal correction. Without rounding, each residual map is $Y \mapsto (3Y - YY^\top Y)/2$, the cubic Newton–Schulz polar iteration in the Björck–Bowie family. These primitives themselves are not new. Before the global cap, the damping weights the undamped reciprocal-gap formula by $\frac{1}{2}h_{ij}^2/(h_{ij}^2 + 0.01^2)$. It therefore suppresses motion relative to that formula when coordinate-diagonal gaps are much smaller than 1% of s; at equal diagonals the motion vanishes. These diagonals need not be eigenvalues, and we did not save spectra or count how many pairs lie in this regime. The statement does not bound absolute motion.

3.3. QR repair and history transport

At positive ages divisible by 10, we refresh both factors. Ages divisible by 50 use QR repair; the remaining refresh ages use Section 3.2. For QR, cast stored C, Q to FP32 and reuse one product:

$$Z = \mathrm{float}(C)\,\mathrm{float}(Q), \quad \pi = \mathrm{argsort}_\downarrow\left(\sum_i \mathrm{float}(Q)_{ij} Z_{ij}\right), \quad Q^+ = R_b\big(\mathrm{qr}(Z_{:,\pi}).Q\big).$$

In exact arithmetic the sorting score equals $\mathrm{diag}(Q^\top C Q)$, even for nonorthogonal Q. Its finite-precision evaluation can change near-tie orderings. The existing reference instead materializes $Q^\top C Q$ for the diagonal and separately forms $CQ_{:,\pi}$, using three dense products rather than one. QR repair is one warm-started orthogonal power step, not an exact eigendecomposition. Over 10,000 calls the candidate performs 9,999 matrix updates with 800 warm and 199 QR refreshes; the reference performs 999 QR refreshes. Auxiliary AdamW performs 10,000 updates in either arm.

After either refresh, transport the newly rounded first moment by reconstructing its world-coordinate representation and projecting into the new bases:

$$M^{\mathrm{world}} = \mathcal{B}(\mathcal{B}(Q_L, M), Q_R^\top), \qquad M^+ = \mathcal{B}(\mathcal{B}((Q_L^+)^\top, M^{\mathrm{world}}), Q_R^+).$$

Here the unprimed bases are the old bases and M denotes $R_b(M_f)$. QR repair permutes $R_b(V_f)$ along both axes according to the sorting orders. Warm refresh leaves that rounded diagonal history unchanged. Neither choice represents an exact rotation of a full second-moment tensor. Eschenhagen et al. report that a basis-aware second-moment correction hurt in their EShampoo experiments (Appendix D.2); this motivates care with such corrections but does not validate our history approximation.

Transpose-based first-moment transport is also approximate when stored bases are nonorthogonal. QR repair has a substantive recovery role: with $Q = I$ and $C = \left[\begin{smallmatrix} 1 & 0.5 \\ 0.5 & 1 \end{smallmatrix}\right]$, all $h_{ij} = 0$, so the warm step leaves the basis unchanged despite nonzero off-diagonal covariance. This counterexample precludes a claim that warm motion alone always diagonalizes a factor; it does not establish incidence in training.

3.4. Compilation and storage

Equal-shaped parameter matrices are batched, retaining independent covariance, basis, and moment tensors. Four groups, each containing 12 matrices, cover shapes 2304×768, 768×768, 3072×768, and 768×1536. Both sides of every matrix are preconditioned: 48 matrices and 96 factors in total. There is no blocking, one-sided approximation, or coverage change relative to the reference.

Separate full-graph compiled regions for ordinary, warm, and QR events include packing parameters/gradients, the matrix update, covariance accumulation, refresh/transport, and parameter writes. Initial eigendecompositions are eager. The measured candidate captures 12 graphs without graph breaks. The compiler uses `emulate_precision_casts=True`. The saved small synthetic CUDA parity probe reports zero maximum parameter difference and zero relative error in all six states for ordinary, warm and QR events. This supports the requested rounding policy on that probe, not an equality proof for every training state. FP32 packing buffers and solver workspaces remain possible.

Per matrix, the six persistent tensors contain $2mn+2m^2+2n^2$ elements. BF16 halves their tensor-byte contribution relative to FP32, while FP32 parameters and auxiliary optimizer state remain. Total GPU peak allocation falls by a smaller fraction than matrix-state storage in these measurements (Section 5.2).

4. Experimental design

We use a 109,396,992-parameter Transformer: width 768, 12 layers, 12 attention heads, feed-forward hidden width 1,536, fused QKV and SwiGLU projections, tied embeddings, RoPE, QK normalization, and no bias or dropout. Model forward uses BF16 autocast with FlashAttention; the model is eager and FP32 cross-entropy is compiled. Every experiment uses one RTX 5090, PyTorch 2.11.0+cu128, and CUDA 12.8 as recorded by the workers.

The cached OpenWebText corpus uses pinned dataset and GPT-2 tokenizer revisions. Documents are split before tokenization/packing; EOS-separated packed rows use 1,024 tokens with causal attention across document boundaries. The experiment selects 320,000 training rows and 1,024 validation rows without replacement, then fixes their order. A global batch contains 32 rows, implemented as four microbatches of eight; each row contributes 1,023 prediction targets. A full run processes 327,360,000 targets over 10,000 calls. Seeds 204 and 205 change model initialization, not data order. Evaluation uses 1,047,552 targets.

The matrix peak learning rate is 0.0015, epsilon is 10^{-8}, weight decay is 0.05 for the declared decayed groups, and global gradient norm is clipped at 1. Auxiliary AdamW uses betas (0.9, 0.999), epsilon 10^{-8}, peak learning rate 0.0006, and the same declared grouping. Both rates use 200 calls of linear warmup followed by cosine decay to 10% of peak at call 10,000. Learning-rate selection and variant development used this research validation panel; it is not an untouched test set. The 0.0015 matrix peak was selected on the seed-204 stock QR10 baseline: three 500-call calibrations at 0.0005, 0.0015 and 0.003 used the full 10k schedule, selecting the lowest call-500 validation CE. That rate was then shared with the warm candidates, without candidate-specific LR retuning.

The primary control is `ResearchSOAP` with FP32 matrix state and QR10. Its update ordering, bias/epsilon convention, QR sorting, and variance permutation are checked against a pinned upstream implementation. Here **stock** denotes the matrix memory settings (0.95, 0.95, covariance 0.95), rather than a claim to reproduce the original SOAP paper's complete training recipe. These values are explicit campaign settings, overriding the local reference constructor's different defaults.

Seed 205 runs reference and candidate sequentially on one claimed GPU. Seed 204 reuses its completed full reference for learning quality and first runs a fresh 201-call stock replay. The replay must reproduce initial weights,

gradients, loss/norm, and validation at calls 0, 50, 100, and 200 exactly. Each candidate repeats its initial backward and matches the paired startup. No model or optimizer tensor checkpoints are reused or saved.

All raw results use frozen source/input hashes and record evaluations, refresh-event counts, finite state, nonnegative variance, and graph captures. Numerical observations cover all factors and physical moments at selected early and late refresh ages. Evaluation checks preserve listed tensors, RNG, and gradient identities; the tensor hash is not a claim to cover every Python scalar or model buffer. Finite moment/Gram measurements are diagnostics rather than validated loss-quality acceptance thresholds.

5. Results

5.1. Learning over the declared horizon

Full validation at call 10,000; two model-initialization seeds, fixed data order.

Seed	Stock QR10 CE	Full BF16 CE	CE increase	Perplexity increase
204	3.411152	3.422389	+0.011237	+1.130%
205	3.409272	3.419354	+0.010082	+1.013%

The mean paired endpoint increase is **0.010659815 CE**. Both seeds show a larger early gap: at call 500, BF16 exceeds the reference by 0.123913 and 0.118438 CE for seeds 204 and 205. From call 2000 onward the observed gaps remain approximately 0.007–0.012 CE, rather than continuing toward zero. The pilot's FP32-state warm arm already shows a comparable call-500 gap (0.118346 CE); all four warm-policy pilot arms span 0.118–0.125 CE there. Thus the early deviation is not specific to BF16 storage. QR20 roughly halves it on this prefix, consistent with a refresh-policy contribution. The pilot does not causally separate cadence, tracking, arithmetic and storage effects at the final horizon. This behavior supports reporting the full curves rather than only endpoint agreement. Two initialization seeds with fixed data order provide limited replication; we make no statistical equivalence claim or confidence-interval interpretation.

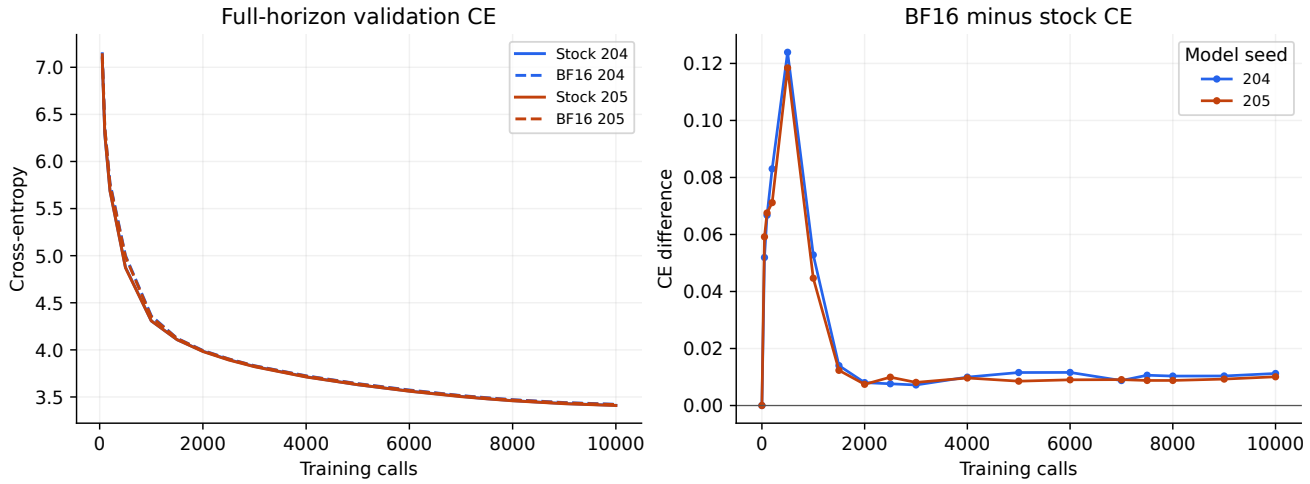


Figure 1. Left: the declared learning horizon, displaying observations from call 50 onward. Right: BF16 minus stock CE, including the shared initial value. Lines connect observed evaluations; they do not represent additional measurements.

5.2. Time and memory

Synchronized steady-step wall means; evaluation, diagnostics and startup excluded.

Seed	Matched calls	Stock step ms	BF16 step ms	Time reduction
204	101–200	380.740	314.638	17.36%
205	101–10000	380.548	309.533	18.66%

Seed 205's timing includes 9,900 calls, covering 198 complete 50-call cycles. Seed 204's newly paired timing includes only calls 101–200 on the replay device; its cached full reference is used for quality, not relabeled as newly paired

timing. These different windows should not be averaged into one speed estimate. The short seed-204 window is within LR warmup but covers two complete 50-call event cycles. Sequential arms share a claimed device; GPU clocks and thermal state were not controlled or logged.

Persistent matrix state falls from 2,491,416,576 to 1,245,708,320 bytes (2.320 to 1.160 GiB; approximately 50%). Seed 205's peak allocated tensor memory falls from 9.397 to 8.223 GiB. Peak allocation includes diagnostic allocations and is distinct from device reservation or persistent state alone. That is approximately 12.49% lower peak allocation, compared with the approximately 50% matrix-state reduction. The BF16 state total includes 32 bytes for four groups' FP32 age/LR scalars.

Complete-step means include forward/backward, matrix packing and writes, and auxiliary AdamW. Diagnostic and evaluation work is outside the timed step scopes and inside campaign elapsed accounting. Compiler caches are shared sequentially within campaigns. The step means do not estimate first-use compilation cost or whole-run time-to-quality, and the combined implementation comparison does not isolate warm tracking as the cause of all time savings. Each forward/backward and optimizer interval is measured with synchronized host wall time, with GPU synchronization at phase boundaries. This prevents overlap across those boundaries; its effect relative to an asynchronous training loop, including any unequal effect on eager and compiled execution, was not isolated.

Seed 205's recorded optimizer means can be accounted for by event class. Mean optimizer time falls from 95.23 to 26.08 ms/call, a 3.65-fold ratio. Ordinary events cost 34.44 ms in stock versus 10.14 ms in the candidate; QR events cost 642.37 versus 601.18 ms, and candidate warm events cost 61.61 ms. The shared timing window contains 90% ordinary events in both arms; stock has 10% QR events, while the candidate has 8% warm and 2% QR events. Consequently 47.28 of the 69.15 ms/call optimizer saving (68.38%) lies in refresh events, and 21.87 ms/call (31.62%) in ordinary events. The whole-step difference is 71.02 ms/call. This identity is descriptive accounting of observed means, not a causal separation of cadence, batching, compilation, arithmetic, and storage. In particular, both refresh frequency and per-event execution differ.

5.3. Short sparse-QR controls

Seed 204; full 10k LR horizon; timing calls 101–2000. Original FP32 stock and cached warm timing come from separate launches; this is mixed-implementation context.

Policy	CE at 500	CE at 2000	Step ms	Optimizer ms
Stock QR10 (pilot)	4.877601	3.985905	382.773	96.358
Warm10/QR50 (cached)	5.001514	3.993939	314.142	26.415
QR20	4.936852	3.990383	329.641	40.460
QR50	5.018046	3.998679	313.551	22.864

These two fresh controls use the same full-BF16 arithmetic, coverage, efficient QR primitive, seed 204, data, and full 10,000-call learning-rate horizon. QR20 sets both refresh and repair periods to 20; QR50 sets both to 50, producing no warm events. The warm comparison reuses the first 2,000 calls of the completed main candidate. Complete shared timing cycles span calls 101–2000. The stock row is the original FP32-state reference's pilot prefix; it is not a matched full-BF16 QR10 control.

At call 2000, warm is 0.003557 CE behind QR20 and 0.004740 ahead of QR50. Warm has lower CE than QR50 at all seven noninitial validation observations; QR20 has lower CE than warm at six. These repeated panels are not independent seeds. The warm prefix has extra refresh work for a modest quality gain over holding bases between QR50s. Its timing comes from the earlier launch, so the small whole-step difference from QR50 is contextual rather than a fresh paired timing result. On this prefix QR20 reaches lower CE at about 4.93% greater measured step time than cached warm. The timing comparison spans launches and supplies no calibrated noise interval. The study establishes neither QR20 superiority over warm at the full horizon nor the converse; both orderings remain open.

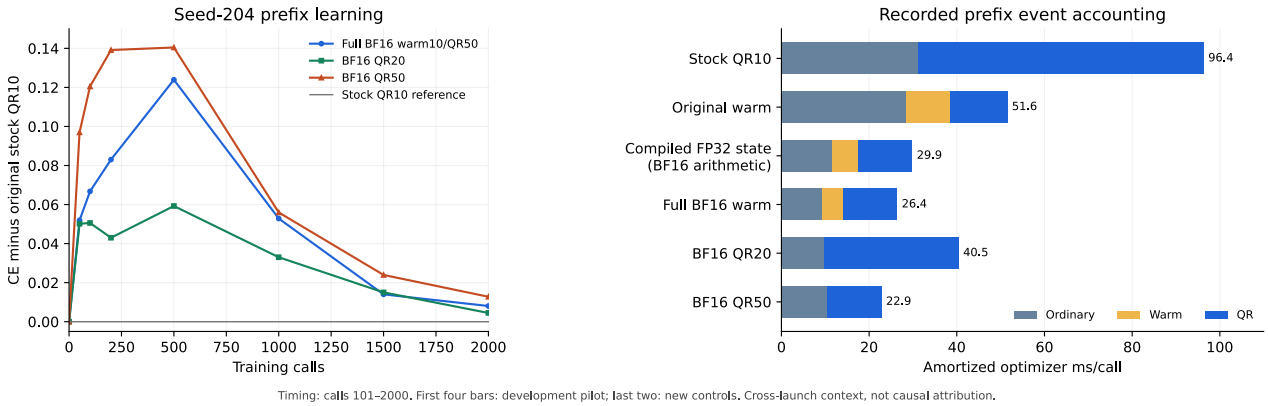


Figure 2. Left: seed204 CE relative to the original stock QR10 prefix. Right: optimizer event-class means weighted by their frequencies over calls 101–2000. The first four bars use the development pilot; QR20/QR50 use the new controls. Original/compiled warm policies use warm10/QR50. The bars describe mixed-implementation, cross-launch context and are not causal ablations.

5.4. What QR product reuse explains

Synthetic matched-input QR primitive; CUDA-event medians, batch 12.

Dimension (batch 12)	Original CUDA ms	Reuse CUDA ms	Time reduction	Reordered columns
768	31.210	30.914	0.95%	0/9216
1536	82.053	80.535	1.85%	12/18432
2304	145.854	141.566	2.94%	40/27648
3072	225.660	215.218	4.63%	160/36864

The profile uses identical fixed synthetic C/Q inputs within each pair, BF16-rounded then promoted to FP32, matched FP32 high matmul precision and FP32 QR output, without compilation or training. Two warmups precede ten counterbalanced measurements per component using CUDA events and synchronized wall time. Batch 12 matches the repeated-shape group size; batch-1 comparisons are retained in the artifact.

The batched primitive saves 0.95–4.63% by reusing CQ. The solver dominates: for dimension 3072, products/sorting fall from 18.165 to 7.724 ms, while the same-input QR solver takes 207.613 ms. Component medians are informative, not an additive decomposition of separately timed code.

Near ties change some column orders, reaching 160 of 36,864 columns in the largest batch; the reordered score band is at most 2.68×10^{-5} relative to the largest score. A synthetic one-factor BF16 transport and next ordinary-update probe, with a fixed 32-dimensional right identity, records a maximum relative difference of 0.0906% between the BF16-rounded reconstructed world moments of the original and reuse probes, and next-update difference of 0.1116%, with cosine at least 0.9999994. These relative differences use the original output's Frobenius norm, floored at 10^{-30}, as denominator. The world-moment reconstruction itself uses BF16 output, unlike the FP32 reconstruction in Section 5.5. Raw stored-coordinate differences can be much larger under signs and permutations. These observations qualify the local comparison; they do not imply bitwise equality or trained-trajectory agreement. We have not trained a new fully optimized FP32-QR10 comparator.

5.5. Numerical observations

For a stored $d \times d$ basis, normalized Gram error is

$$e_Q = \frac{\|Q^\top Q - I\|_F}{\sqrt{d}}.$$

For each sampled matrix refresh, let

$$P_{\text{before}} = Q_L R_b(M_f) Q_R^\top, \qquad P_{\text{after}} = Q_L^+ M^+ (Q_R^+)^\top, \qquad e_M = \frac{\|P_{\text{after}} - P_{\text{before}}\|_F}{\max(\|P_{\text{before}}\|_F, 10^{-12})}.$$

All stored tensors in these diagnostic reconstructions are promoted to FP32 and their products use "highest" matmul precision. The reference is the candidate's newly rounded first moment in its old bases immediately before transport. Thus e_M is a local before/after refresh discrepancy, rather than deviation from stock SOAP or total accumulated history error. Reported maxima range over sampled factors or matrices and observation calls.

All sampled main-run states remain finite and candidate captures stay at 12 graphs without graph breaks. Maximum observed physical-M transport error is 0.9924% for seed 204 and 1.0356% for seed 205. Maximum normalized stored-basis Gram error is approximately 0.002366 and 0.002367. The historical record retains the seed-205 crossing of the 1% M screen and both candidates' crossings of the original absolute Gram screen. These engineering screens are not calibrated predictors of loss quality and do not define method acceptance here. The seed205 1% moment crossing occurs only at the first sampled warm refresh (age 10); all later sampled maxima are at most 0.5424%.

The saved coordinate residual directly measures diagonalization:

$$r_C = \frac{\|\operatorname{offdiag}(Q^\top C Q)\|_F}{\max(\|Q^\top C Q\|_F, 10^{-30})},$$

evaluated with FP32 products at "highest". At seed205 ages 1990 and 9990, after four warm refreshes since the preceding QR, the candidate's unweighted means across 96 factors are 0.394 and 0.380, versus 0.201 and 0.214 for QR10. Candidate maxima are 0.965 and 0.975. This describes poorer covariance diagonalization between repairs, despite small Gram errors; Frobenius weighting can obscure weaker spectral directions.

Rounding the observed dense bases to BF16 introduces measurable orthogonality error even after a fresh FP32 QR; this is an empirical rounding baseline, not a universal lower bound. Normalized Gram and actual-M observations describe different aspects of the state. Neither alone certifies diagonalization, adaptive-direction fidelity, or convergence. The equal-diagonal stall and approximate diagonal-V history are therefore retained as limitations even when the sampled state remains finite.

6. Discussion and limitations

The measured approximation follows the reference with a modest final gap on this workload, halves matrix-state storage, and lowers measured step time. That result is useful as an implementation operating point. It does not show that warm tracking, BF16 storage, or the particular repair schedule is generally optimal. Storage and compute changes also alter numerical trajectories.

The evidence covers one model scale, one dataset cohort/order, two model initializations, one GPU family, and a validation panel used during development. There are no generalization claims across architectures, corpora, or hardware, and no untouched test result or statistical equivalence test. Short QR20/QR50 controls cannot determine a final-horizon ranking. Synthetic QR inputs and one-factor probes do not represent the trained distribution of all factor pairs.

The primary time comparison uses the original audited FP32 QR10 implementation. Our isolated reuse profile narrows one implementation attribution question but does not replace a newly trained efficient-reference comparison. There is no direct matched evaluation against subspace or adaptive-refresh methods, and the warm step has no standalone diagonalization guarantee.

The early-gap pattern is qualitatively consistent with Eschenhagen et al.'s finding that accurate bases mattered more early in their Imagewoof ViT experiments (§4.4). That different workload does not establish the mechanism of our gap. Our fixed schedule and reference-calibrated LR leave that question open. Ignoring startup and other excluded work, the seed205 mean-step ratio assigns 10,000 candidate calls the same measured steady cost as approximately 8,134 stock calls. Whether a reference with an appropriately shortened LR horizon would reach the same loss is unmeasured. This arithmetic is not a time-to-quality result.

7. Conclusion

TurboSOAP specifies and evaluates a dense BF16-state SOAP approximation with compiled matrix updates, damped tracking, and periodic QR repair. The two-seed case study yields a small consistent endpoint loss increase with lower matrix state and measured step time against the declared reference. Explicit precision and history rules, local QR attribution, and short cadence controls delimit that tradeoff and support reproducible follow-up comparisons.

Appendix A. Implementation-family context

An earlier five-arm 2,000-call pilot used the same seed-204 recipe and a newly paired timing campaign. It included original stock QR10, original FP32-state warm, a compiled FP32-storage mode with BF16 matrix arithmetic, full BF16, and BF16 moments/covariance with FP32 bases.

Development pilot, seed 204; call-2000 CE and calls 101–2000 timing on the full 10k LR horizon; not an additional independent seed.

Development arm	CE at 2000	Step ms	Matrix-state GiB
Stock QR10	3.985905	382.773	2.320
Original warm, FP32 state	3.995753	337.172	2.320
Compiled FP32 storage / BF16 arithmetic	3.991560	316.281	2.320
Full BF16	3.993939	313.162	1.160
BF16 M/V/C, FP32 bases	3.995611	315.376	1.608

The compiled FP32-storage mode's name describes storage, not all-FP32 arithmetic. Switching to full BF16 changes stored bases and their associated correction arithmetic as well as moment/covariance storage. Its roughly 0.99% lower step time in this pilot is consequently a within-family option comparison, not a pure storage-only causal attribution. The pilot is development evidence, not additional independent seed replication; its shared prefix must not be counted as a new independent full-BF16 observation. The full-BF16 endpoint is 0.002379 CE above the compiled FP32-storage arm. Because basis precision and correction arithmetic also differ, that contrast does not identify a pure storage penalty or partition the final-horizon gap.

Appendix B. Reproducibility pointers

The repository provides the matrix-only `bf16_soap.CompiledSOAP` entry point, frozen campaign manifests/protocols, scalar logs and audit receipts. The reproducibility guide (`REPRODUCIBILITY.md`) identifies exact configurations, software, data revisions and hashes, CPU-only paper regeneration, and the historical launch wrappers. The claim map (`CLAIMS.md`) and machine-readable evidence inventory connect tables and figures to their source artifacts. Experiments save no trained model/optimizer tensor checkpoints. All reported figures and numeric tables can be reconstructed from the scalar evidence.

References

Nikhil Vyas, Depen Morwani, Rosie Zhao, Mujin Kwun, Itai Shapira, David Brandfonbrener, Lucas Janson, and Sham Kakade. *SOAP: Improving and Stabilizing Shampoo using Adam.* arXiv:2409.11321, 2024; v2, 2025. Paper.

Alan Milligan, Zikun Xu, Simon Lacoste-Julien, Felix Dangel, and Wu Lin. *Reparametrizing Shampoo and SOAP for Subspace Basis Updates and BFloat16 Storage.* arXiv:2605.26327v1, 2026. Paper.

Runa Eschenhagen, Aaron Defazio, Tsung-Hsien Lee, Richard E. Turner, and Hao-Jun Michael Shi. *Purifying Shampoo: Investigating Shampoo's Heuristics by Decomposing its Preconditioner.* arXiv:2506.03595v2, 2025. Paper.

G. E. Forsythe and P. Henrici. *The Cyclic Jacobi Method for Computing the Principal Values of a Complex Matrix.* Transactions of the American Mathematical Society 94(1):1–23, 1960. Paper.

Å. Björck and C. Bowie. *An Iterative Algorithm for Computing the Best Estimate of an Orthogonal Matrix.* SIAM Journal on Numerical Analysis 8(2):358–364, 1971. Paper.

Sike Wang, Pan Zhou, Jia Li, and Hua Huang. *4-bit Shampoo for Memory-Efficient Network Training.* arXiv:2405.18144, 2024; v3, 2025. Paper.